


Problem Set 3 Image Mosaics

In this exercise, you will implement an image stitcher that uses image warping and homographies to automatically create an image mosaic. We will focus on the case where we have two input images that should form the mosaic, where we warp one image into the plane of the second image and display the combined views. This problem will give some practice manipulating homogeneous coordinates, computing homography matrices, and performing image warps. For simplicity, we will specify corresponding pairs of points manually using mouse clicks.

Figure 1: Image Mosaics



Here are the main components you will need to implement:

1 Getting correspondences

Use the provided code to get manually identified corresponding points from two views. “ginput” function provides an easy way to collect mouse click positions. The results will be sensitive to the accuracy of the corresponding points, when providing clicks, choose distinctive points in the image that appear in both views.

2 Computing the homography parameters

Write a function that takes a set of corresponding image points and computes the associated 3×3 homography matrix H . This matrix transforms any point p in one view to its corresponding homogeneous coordinates in the second view, p' , such that $p' = Hp$. Note that p and p' are both $3D$ points in homogeneous coordinates. The function should take a list of $n - 4$ pairs of corresponding points from the two views, where each point is specified with its $2D$ image coordinates.



We can set up a solution using a system of linear equations $Ax = b$, where the 8 unknowns of H are stacked into an 8-vector x , the $2n$ -vector b contains image points from one view, and the $2n \times 8$ matrix A is filled appropriately so that the full system gives us $\lambda p' = Hp$. There are only 8 unknowns in H because we set $H_{33} = 1$. Solve for the unknown homography matrix parameters. [Useful numpy functions: 'lstsq', reshape]

Verify that the homography matrix your function computes is correct by mapping the clicked image points from one view to the other, and displaying them on top of each respective image (matplotlib: imshow, followed by hold and plot). Be sure to handle homogenous and non-homogenous coordinates correctly.

3 Warping between image planes

Write a function that can take the recovered homography matrix and an image, and return a new image that is the warp of the input image using H . Since the transformed coordinates will typically be sub-pixel values, you will need to sample the pixel values from nearby pixels. For color images, warp each RGB channel separately and then stack together to form the output.

To avoid holes in the output, use an inverse warp. Warp the points from the source image into the reference frame of the destination, and compute the bounding box in that new reference frame. Then sample all points in that destination bounding box from the proper coordinates in the source image (linear interpolation). Note that transforming all the points will generate an image of a different shape / dimensions than the original input.

[Useful functions: numpy.meshgrid, scipy.interpolate.RectBivariateSpline]

4 Create the output mosaic

Once we have the source image warped into the destination images frame of reference, we can create a merged image showing the mosaic. Create a new image large enough to hold both (registered) views; overlay one view onto the other, simply leaving it black wherever no data is available. Do not worry about artifacts that result at the boundaries.

5 Requirements

After writing and debugging your system:

1. Apply your system to the provided pair of images, and display the output mosaic.

2. Show one additional example of a mosaic you create using images that you have taken. You might make a mosaic from two or more images of a broad scene that requires a wide angle view to see well. Or, make a mosaic using two images from the same room where the same person appears in both.
3. [Optional] Warp one image into a frame region in the second image. To do this, let the points from the one view be the corners of the image you want to insert in the frame, and let the corresponding points in the second view be the clicked points of the frame (quadrilateral) into which the first image should be warped. Use this idea to replace one surface in an image with an image of something else. For example – overwrite a billboard with a picture of your choice, or project a drawing from one image onto the street in another image, or replace a portrait on the wall with someone else's face, or paste a Powerpoint slide onto a movie screen,.

For all examples, play around a bit with the choice of points for the correspondence pairs until you get a reasonable alignment. Be sure to comment your code.

Figure 2: Image Mosaics



6 Extra credit [up to 10 pts each]

1. Replace the manual correspondence stage with automatic interest point detection and local feature matching. Check out available code and OpenCV functions here to compute the local interest points and features:
 - http://docs.opencv.org/modules/features2d/doc/feature_detection_and_description.html
 - <http://www.vlfeat.org/overview/sift.html>
 - <http://www.robots.ox.ac.uk/~vgg/research/affine/detectors.html>



2. Implement RANSAC for robustly estimating the homography matrix from noisy correspondences. Show with an example where it successfully gives good results even when there are outlier (bad) correspondences given as input. Compare the robust output to the original (non-RANSAC) implementation where all correspondences are used.

7 Deliverables

You are required to deliver the following:

- Your code.
- Output for some test images.
- Report including explanation of your code and representative results on sample test images.

8 References

- Course lectures 10 and 11.
- The original matlab-based assignment can be found at <http://www.cs.utexas.edu/~grauman/courses/spring2011/psets/pset3/>. Sample students' results can be found at <http://vision.cs.utexas.edu/cv-spring2011/mosaics/>.

Good Luck